



# Machine learning

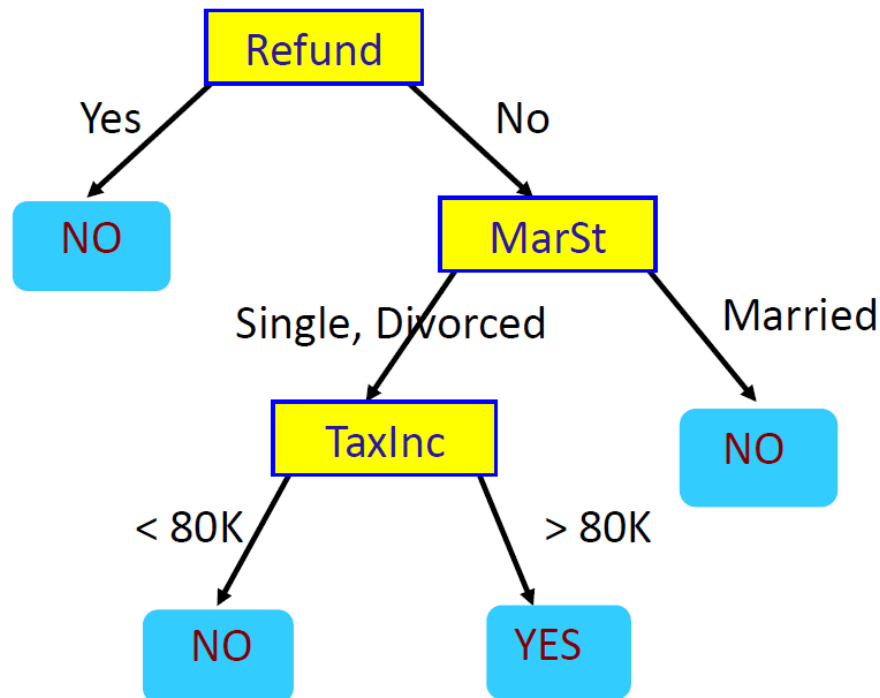
## Decision Trees

Mohammad-Reza A. Dehaqani

[dehaqani@ut.ac.ir](mailto:dehaqani@ut.ac.ir)

Slides are mainly adopted from cmu Aarti course

# Decision Trees; discrete features, tax fraud detection



$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat

- Each internal node: test one feature  $X_i$
- Each branch from a node: selects some value for  $X_i$
- Each leaf node: prediction for  $Y$

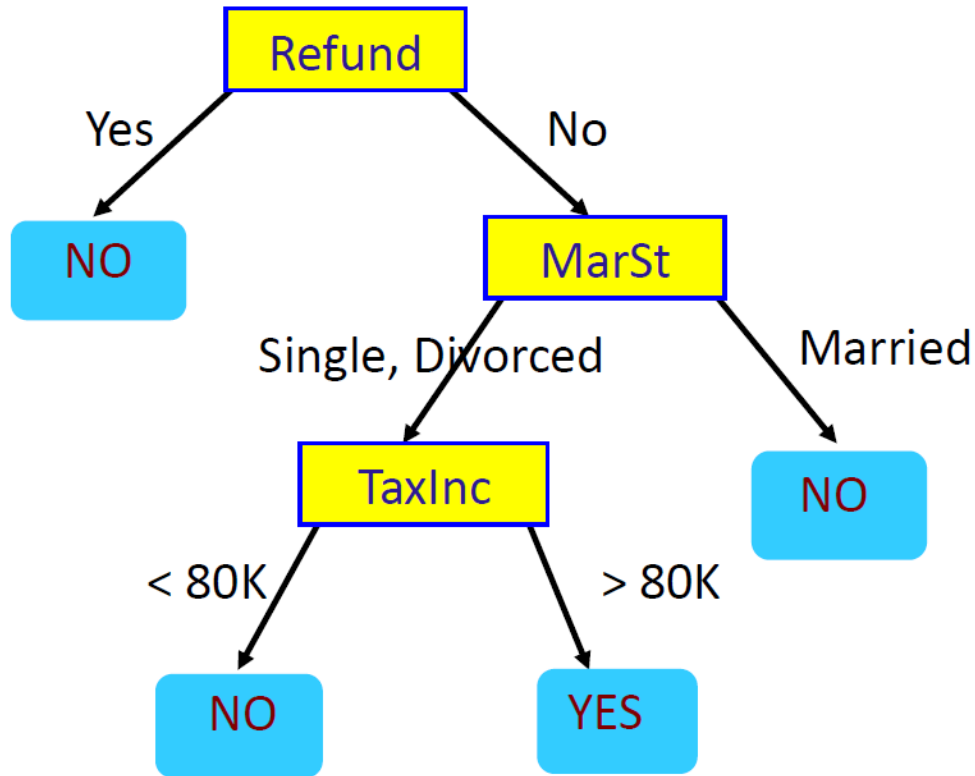
**Prediction:** Given a decision tree, how do we assign label to a test point

# Decision Tree for Tax Fraud Detection

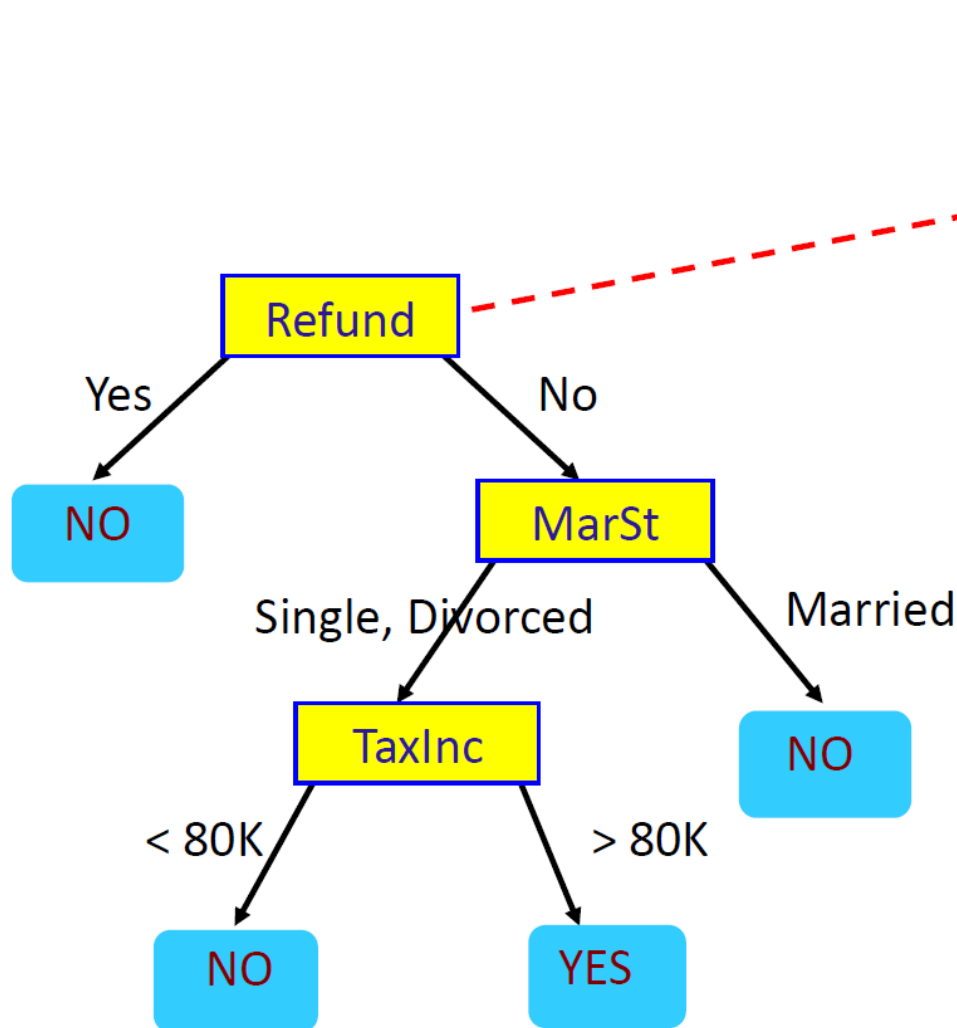


Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



# Decision Tree for Tax Fraud Detection



Query Data

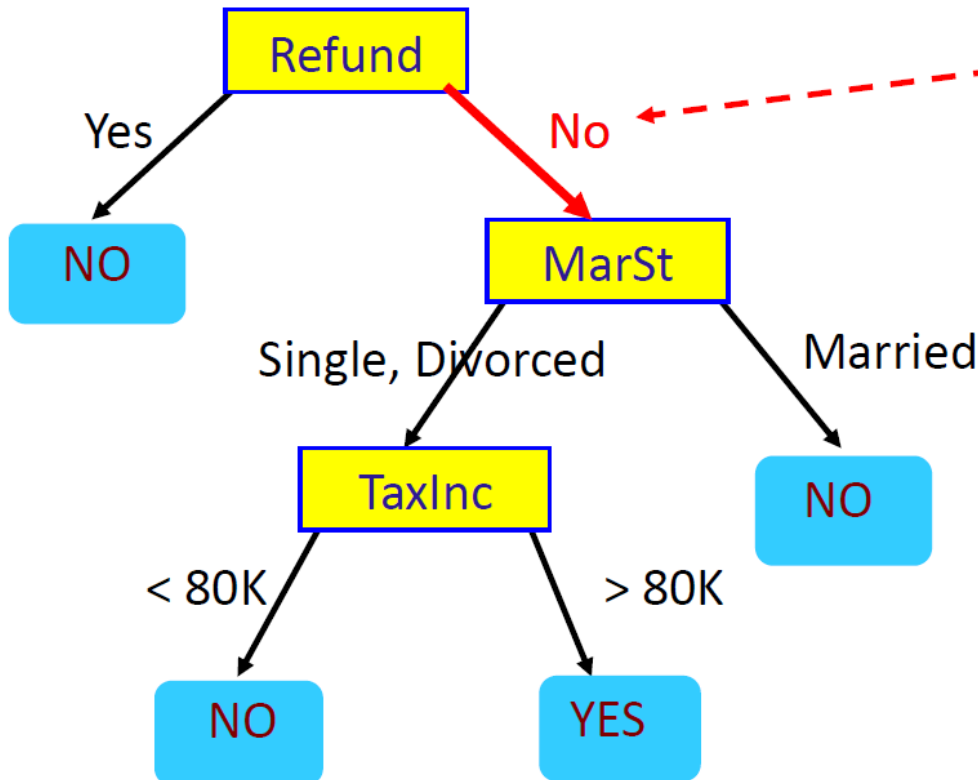
$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

# Decision Tree for Tax Fraud Detection

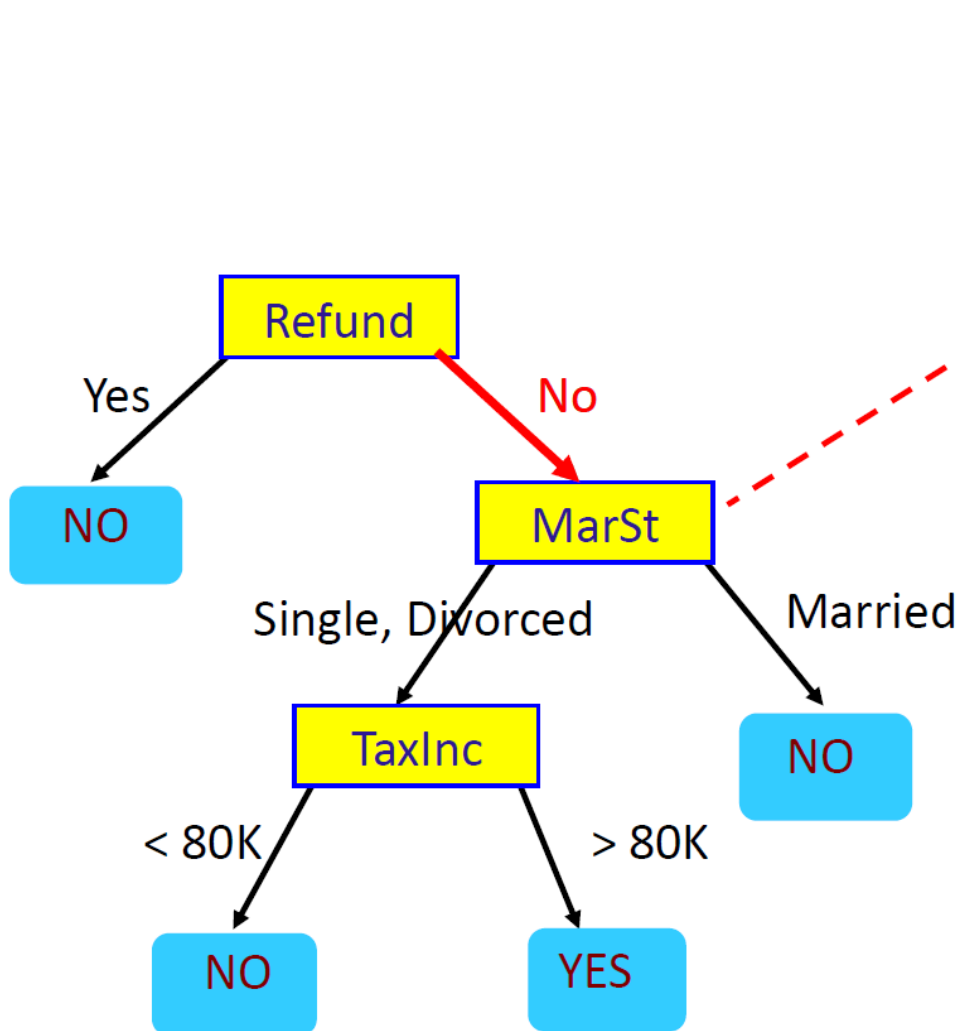


Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



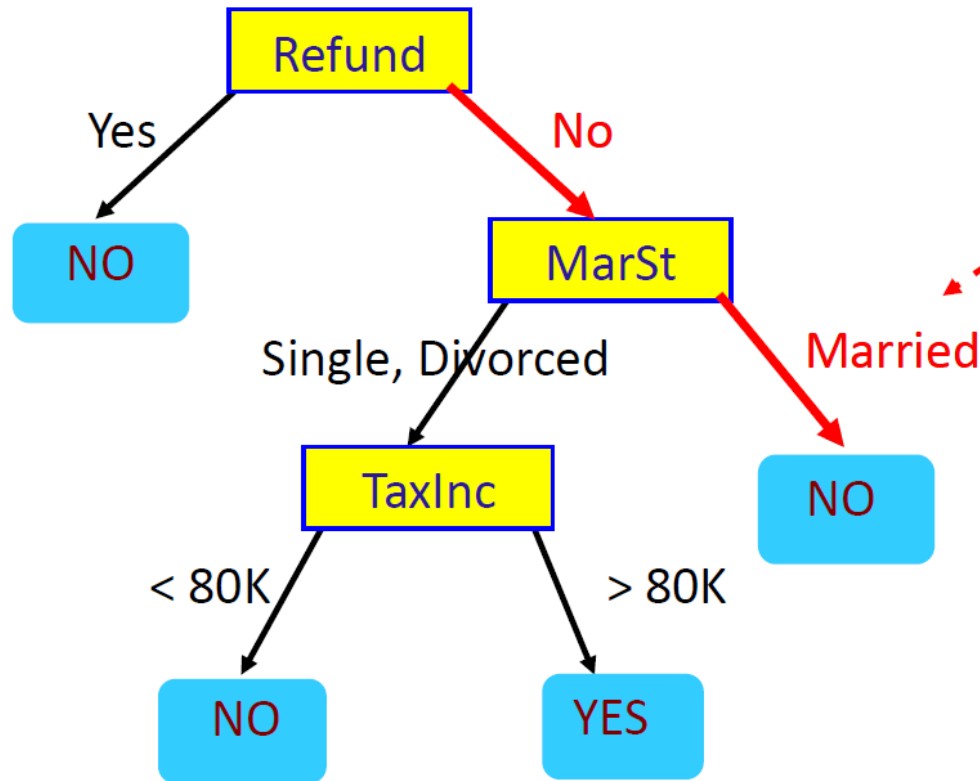
# Decision Tree for Tax Fraud Detection



Query Data


$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

# Decision Tree for Tax Fraud Detection



Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

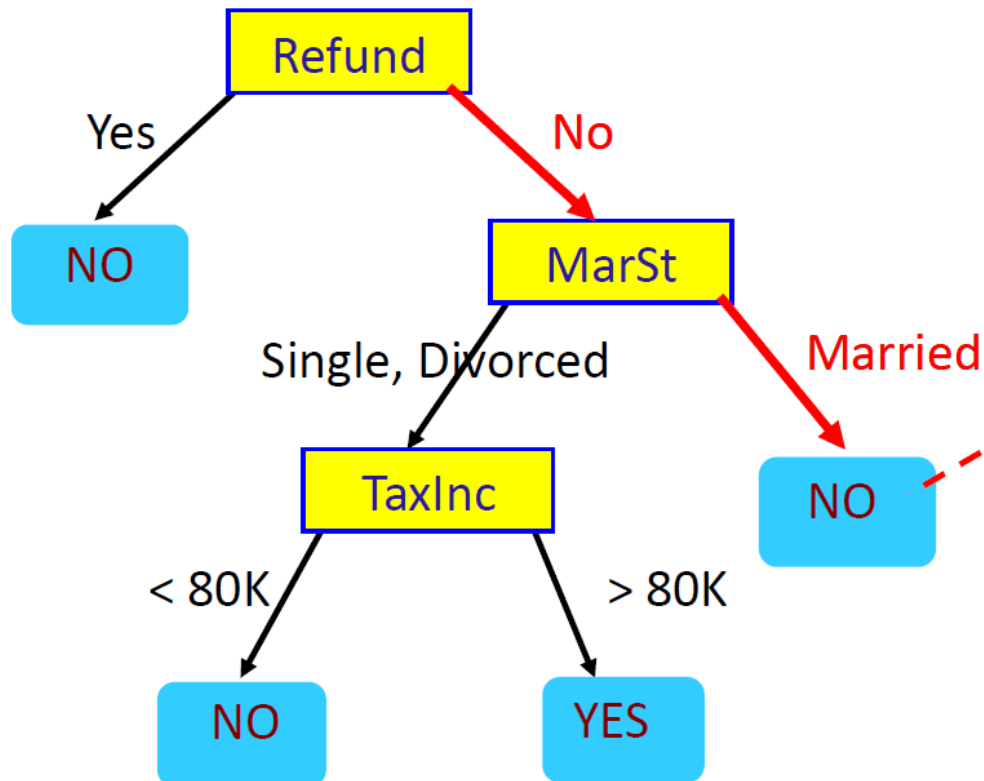


# Decision Tree for Tax Fraud Detection



Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Assign Cheat to "No"



# So far...



- What does a decision tree represent
- Given a decision tree, how do we assign label to a test point

Discriminative or Generative?

**Now ...**

- How do we learn a decision tree from training data

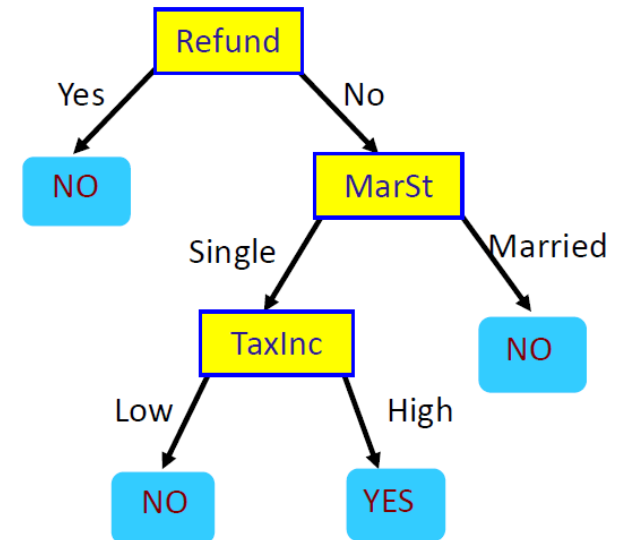
# How to learn a decision tree



- Top-down induction [ID3]

Main loop:


1.  $X \leftarrow$  the “best” decision feature for next *node*
2. Assign  $X$  as decision feature for *node*
3. For each value of  $X$ , create new descendant of *node* (Discrete features)
4. Sort training examples to leaf nodes
5. If training examples perfectly classified, Then STOP, Else iterate over new leaf nodes (steps 1-5) after removing current feature
6. When all features exhausted, assign majority label to the leaf node



# Which feature is best?



$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F




Y: 4 Ts  
0 Fs

Y: 1 Ts  
3 Fs

Absolutely  
sure

Kind of  
sure



Y: 3 Ts  
1 Fs



Y: 2 Ts  
2 Fs

Kind of  
sure

Absolutely  
unsure

Good split if we are more certain  
about classification after split –  
Uniform distribution of labels is bad

# Which feature is best?



Pick the attribute/feature which yields maximum information gain:

$$\arg \max_i I(Y, X_i) = \arg \max_i [H(Y) - H(Y|X_i)]$$

$H(Y)$  – entropy of  $Y$       $H(Y|X_i)$  – conditional entropy of  $Y$

# Entropy




- Entropy of a random variable  $Y$

$$H(Y) = - \sum_y P(Y = y) \log_2 P(Y = y)$$

***More uncertainty,  
more entropy!***

$Y \sim \text{Bernoulli}(p)$



**Information Theory interpretation:**  $H(Y)$  is the expected number of bits needed to encode a randomly drawn value of  $Y$  (under most efficient code)

# Information Gain



- Advantage of attribute = decrease in uncertainty
  - Entropy of  $Y$  before split

$$H(Y) = - \sum_y P(Y = y) \log_2 P(Y = y)$$

- Entropy of  $Y$  after splitting based on  $X_i$ 
  - Weight by probability of following each branch

$$\begin{aligned} H(Y | X_i) &= \sum_x P(X_i = x) H(Y | X_i = x) \\ &= - \sum_x P(X_i = x) \sum_y P(Y = y | X_i = x) \log_2 P(Y = y | X_i = x) \end{aligned}$$

- Information gain is difference

$$I(Y, X_i) = H(Y) - H(Y | X_i)$$

**Max Information gain = min conditional entropy**

# Which feature is best to split?



Pick the attribute/feature which yields maximum information gain:

$$\begin{aligned}\arg \max_i I(Y, X_i) &= \arg \max_i [H(Y) - H(Y|X_i)] \\ &= \arg \min_i H(Y|X_i)\end{aligned}$$

Entropy of Y

$$H(Y) = - \sum_y P(Y = y) \log_2 P(Y = y)$$

Conditional entropy of Y

$$H(Y | X_i) = \sum_x P(X_i = x) H(Y | X_i = x)$$



Feature which yields maximum reduction in entropy (uncertainty) provides maximum information about Y

# Information Gain



$$H(Y | X_i) = - \sum_x P(X_i = x) \sum_y P(Y = y | X_i = x) \log_2 P(Y = y | X_i = x)$$

$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F



$$\hat{H}(Y|X_1) = -\frac{1}{2}[1 \log_2 1 + 0 \log_2 0] - \frac{1}{2}\left[\frac{1}{4} \log_2 \frac{1}{4} + \frac{3}{4} \log_2 \frac{3}{4}\right]$$

$$\hat{H}(Y|X_2) = -\frac{1}{2}\left[\frac{3}{4} \log_2 \frac{3}{4} + \frac{1}{4} \log_2 \frac{1}{4}\right] - \underbrace{\frac{1}{2}\left[\frac{1}{2} \log_2 \frac{1}{2} + \frac{1}{2} \log_2 \frac{1}{2}\right]}_{> 0}$$

$$\hat{H}(Y|X_1) < \hat{H}(Y|X_2)$$



# Handling continuous features



Convert continuous features into discrete by setting a threshold.



What threshold to pick?

Search for best one as per information gain. Infinitely many??

Don't need to search over more than  $\sim n$  (number of training data), e.g. say  $X_1$  takes values  $x_1^{(1)}, x_1^{(2)}, \dots, x_1^{(n)}$  in the training set. Then possible thresholds are

$$[x_1^{(1)} + x_1^{(2)}]/2, [x_1^{(2)} + x_1^{(3)}]/2, \dots, [x_1^{(n-1)} + x_1^{(n)}]/2$$



# Dyadic decision trees (split on mid-points of features)



# Decision Tree more generally



$$X_1 \geq 0.5, X_2 = \{a, b\} \text{ or } \{c, d\}$$




- Features can be discrete, continuous or categorical
- Each internal node: test some set of features  $\{X_i\}$
- Each branch from a node: selects a set of value for  $\{X_i\}$
- Each leaf node: prediction for Y

# Expressiveness of Decision Trees



- Decision trees in general (without pruning) can express any function of the input features.
- E.g., for Boolean functions, truth table row  $\rightarrow$  path to leaf:



- There is a decision tree which perfectly classifies a training set with one path to leaf for each example - **overfitting**
- But it won't generalize well to new examples - prefer to find more **compact** decision trees

# When to Stop?



- Many strategies for picking simpler trees:

- Pre-pruning


- Fixed depth (e.g. ID3)
- Fixed number of leaves

- Post-pruning

- Chi-square test

- Convert decision tree to a set of rules
- Eliminate variable values in rules which are independent of label (using chi-square test for independence)
- Simplify rule set by eliminating unnecessary rules

- Information Criteria: MDL (Minimum Description Length)



# Information Criteria



- Penalize complex models by introducing cost



$$\hat{f} = \arg \min_T \left\{ \underbrace{\frac{1}{n} \sum_{i=1}^n \text{loss}(\hat{f}_T(X_i), Y_i)}_{\text{log likelihood}} + \underbrace{\text{pen}(T)}_{\text{cost}} \right\}$$

$$\begin{aligned} \text{loss}(\hat{f}_T(X_i), Y_i) &= (\hat{f}_T(X_i) - Y_i)^2 && \text{regression} \\ &= \mathbf{1}_{\hat{f}_T(X_i) \neq Y_i} && \text{classification} \end{aligned}$$

$\text{pen}(T) \propto |T|$       penalize trees with more leaves

CART – optimization can be solved by dynamic programming

# Example of 2-feature decision tree classifier




# How to assign label to each leaf



Classification – Majority vote

Regression – ?







# How to assign label to each leaf



Classification – Majority vote



Regression – Constant/  
Linear/Poly fit

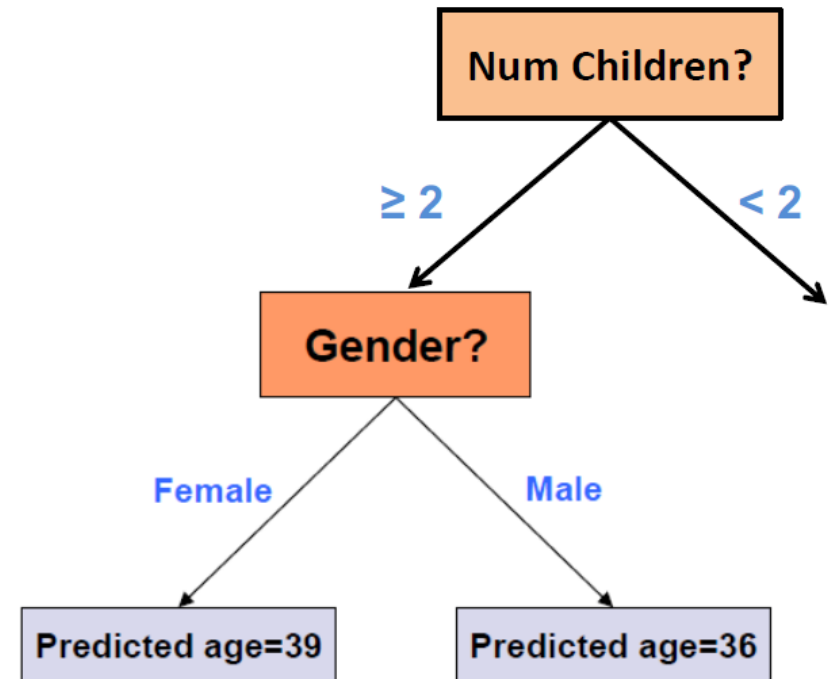


# Regression trees



$X^{(1)}$       ....       $X^{(p)}$      $Y$

Gender	Rich?	Num. Children	# travel per yr.	Age
F	No	2	5	38
M	No	0	2	25
M	Yes	1	0	72
:	:	:	:	:



Average (fit a constant ) using training data at the leaves

# What you should know



- Decision trees are one of the most popular data mining tools
  - Simplicity of design
  - Interpretability
  - Ease of implementation
  - Good performance in practice (for small dimensions)
- Information gain to select attributes (ID3, C4.5,...)
- Decision trees will overfit!!!
  - Must use tricks to find “simple trees”, e.g.,
    - Pre-Pruning: Fixed depth/Fixed number of leaves
    - Post-Pruning: Chi-square test of independence
    - Complexity Penalized/MDL model selection
- Can be used for classification, regression and density estimation too

